
OVERVIEW

Experienced AI Engineer with over 5 years of expertise leveraging data and generative AI to design and implement impactful solutions in investment banking. Skilled in building and optimizing ETL pipelines, fine-tuning large language models, and automating workflows to enable high-quality data-driven insights. Strong emphasis on driving operational efficiency and improving financial reporting through scalable AI and data engineering practices.

SKILLS

- Languages: Python, C#, SQL, SnowQL, **Rust**, C, Bash Linux
- Version Control: Github, Agile Practices, SDLC, Docker, Kubernetes, CI/CD
- Data Engineering : ELT/ ELT, Data Modelling , Machine learning, LLMs
- Databases: Snowflake: Postgres, MS-SQL. MongoDB,

WORK EXPERIENCE

SENIOR TECHNOLOGY ASSOCIATE - DATA & AI - MORGAN STANLEY

June 2025 - Present

1

- Developing Data Infrastructure and models to design semantic models that Generative AI models can use as extensive context outside of reporting.
- Built and Developed a **data quality** matrix for different business components to provide better visibility on data quality and accuracy analysis across several dimensions and business units using **dbt models**, tests and pre and post hooks to establish a clean error lineage.
- Developed **multi-agent systems** for parsing and analyzing massive text documents, enabling scalable and efficient processing of large-scale data.
- Led design and development of ETL pipelines and data models to support key stakeholder reporting and analytics in an investment banking environment, automating data ingestion and validation processes using Python.
- Architected and optimized **Snowflake** ETL pipelines using secure SQL, including writing efficient DDLs, stored procedures, and transformation scripts following best practices such as stepwise **data transformations**, incremental loading, and **data validation** to ensure reliability and maintainability.
- Delivered firm-wide training on best practices for integrating Agentic AI and **GitHub Copilot** into workflows, driving adoption of AI-powered automation and improving operational efficiency across data engineering teams.

AI ENGINEER - INDEPENDENT CONSULTANT

March 2024 - June 2025

- **Independently designed and developed full-stack** generative AI applications in startup environments, managing the complete model lifecycle from training through inference to deployment.
- **Gained deep expertise by fine-tuning, evaluating,** and deploying client-facing large language models (LLMs) using open-source frameworks such as LangChain, LlamaIndex, and Hugging Face, adapting to fast-moving project needs.
- Created Python and Bash automation scripts for seamless deployment, monitoring, and scaling of AI models and infrastructure, **enhancing operational efficiency** in resource-constrained startup settings.
- **Engineered GPU-accelerated training and inference pipelines** optimized for maximum performance and cost-efficiency, supporting rapid iteration cycles and production readiness.
- **Developed domain-specific retrieval-augmented generation (RAG) systems,** integrating extensive data sources to improve the relevance and accuracy of LLM outputs in niche applications.
- **Built inference and prompt engineering frameworks** aligned with ISO quality standards, incorporating validation mechanisms to ensure coherence, compliance, and robustness in generated results.

ANALYTICS ENGINEER - RESONANCE COMPANIES

September 2022 - November 2023

- Data Modeling with **dbt**: Extensive use of dbt for **data modeling** and transformation within Snowflake.
- **Machine Learning with Python**: Developed machine learning models using Python for tasks such as predicting product defects.
- Designed and implemented Snowflake-based data models and pipelines as the central finance data warehouse, integrating high-quality diverse data sources to support accurate revenue and expense reporting and enable timely decision-making.
- **Optimized ETL workflows** including DDLs, stored procedures, and transformation scripts following best practices for secure SQL coding, incremental loading, and performance tuning to ensure scalable, reliable data processing.
- **Developed automated CI/CD pipelines** leveraging Docker, Kubernetes, and Argo CD to deploy and scale data solutions efficiently, reducing deployment errors and manual efforts in data operations.
- **Created actionable insights using Looker dashboards** for finance stakeholders, aligning reporting metrics to business objectives, and driving transparency and process improvements in the financial value chain.

DATA ENGINEER - PENPATH LLC

August 2021 - September 2022

- **Data Warehousing & Pipelines:** Created data warehousing infrastructure and pipelines, leveraging **Snowflake**, to deliver marketing analysis and drive client sales.
- **dbt Model Development:** Extensive experience in dbt model development for data transformation and analysis within Snowflake.
- **Tableau Dashboard Management:** Managed and maintained Tableau dashboards for data visualization and reporting.
- **Multi-Touch Attribution Modeling:** Developed Markov chain models in **Python** for multi-touch attribution marketing, extracting channel-specific revenue impacts.
- **Python-based Data Integration:** Developed data integration pipelines for Python-based machine learning algorithms.
- **ETL Collaboration:** Collaborated with stakeholders and marketing experts to understand the data better, which helped us build better data pipelines and dashboards.

SOFTWARE ENGINEER - LETSCOOEE

July 2020 - August 2021

- Designed and deployed end-to-end mobile applications using **Flutter**, web-based **APIs** with **Node.js**, **Mongoose**, and **MongoDB**.
- Built and deployed templates for HTML-based games for clients.

EDUCATION

BSC IN COMPUTER SCIENCE - UNIVERSITY OF MUMBAI - 2018

First in class - Focus on Fundamentals of Data structures, Databases, Networking

Published Paper on Agentic AI in JETIR journal -

<https://www.jetir.org/papers/JETIRAB06180.pdf>